

Name: Dongming Liu
NetID: dl35
Section: AL

ECE 408/CS483 Milestone 3 Report

0. List Op Times, whole program execution time, and accuracy for batch size of 100, 1k, and 10k images from your basic forward convolution kernel in milestone 2. This will act as your baseline this milestone.

Batch Size	Op Time 1	Op Time 2	Total Execution Time	Accuracy
100	0.175719 ms	0.633705 ms	0m1.220s	0.86
1000	1.59055 ms	6.10214 ms	0m9.758s	0.886
10000	15.7009 ms	60.7863 ms	1m35.228s	0.8714

1. **Optimization 1: Weight matrix in constant memory (1 points)**

- a. Which optimization did you choose to implement? Chose from the optimization below by clicking on the check box and explain why did you choose that optimization technique.

- ☐ Tiled shared memory convolution (2 points)
- ☐ Shared memory matrix multiplication and input matrix unrolling (3 points)
- ☐ Kernel fusion for unrolling and matrix-multiplication (2 points)
- ☒ Weight matrix in constant memory (1 point)
- ☐ Tuning with restrict and loop unrolling (3 points)
- ☐ Sweeping various parameters to find best values (1 point)
- ☐ Multiple kernel implementations for different layer sizes (1 point)
- ☐ Input channel reduction: tree (3 point)
- ☐ Input channel reduction: atomics (2 point)
- ☐ Fixed point (FP16) arithmetic. (4 points)
- ☐ Using Streams to overlap computation with data transfer (4 points)
- ☐ An advanced matrix multiplication algorithm (5 points)
- ☐ Using Tensor Cores to speed up matrix multiplication (5 points)
- ☐ Overlap-Add method for FFT-based convolution (8 points)
- ☐ Other optimizations: please explain

I select weight matrix in constant memory as one of my optimizations. I picked this one because it is very straightforward, and this technique appears in one of the previous MPs.

- b. How does the optimization work? Did you think the optimization would increase performance of the forward convolution? Why? Does the optimization synergize with any of your previous optimizations?

The optimization works by transferring the weight mask to the constant memory and reducing global memory access. Since the convolution kernel never updates the mask, it is better to store the weight mask on constant memory. I thought that this technique would increase the performance of the forward convolution. Constant memory is faster than global memory in general, so using constant memory can reduce global memory access. This is the first optimization.

- c. List the Op Times, whole program execution time, and accuracy for batch size of 100, 1k, and 10k images using this optimization (including any previous optimizations also used).

Batch Size	Op Time 1	Op Time 2	Total Execution Time	Accuracy
100	0.166936	0.570333 ms	0m1.204s	0.86
1000	1.47756 ms	5.6494 ms	0m9.715s	0.886
10000	14.6025 ms	56.799 ms	1m35.093s	0.8714

- d. Was implementing this optimization successful in improving performance? Why or why not? Include profiling results from *nsys* and *Nsight-Compute* to justify your answer, directly comparing to your baseline (or the previous optimization this one is built off of)

The optimization was successful in improving performance.

```

71 Generating CUDA Kernel Statistics...
72
73 Generating CUDA Memory Operation Statistics...
74 CUDA Kernel Statistics (nanoseconds)
75
76 Time(%)    Total Time    Instances    Average    Minimum    Maximum    Name
77 -----
78 100.0      76774225         2    38387112.5    15883199    60891026    conv_forward_kernel
79  0.0         2848         2      1424.0        1408         1440    do_not_remove_this_kernel
80  0.0         2720         2      1360.0        1312         1408    prefn_marker_kernel

```

Kernel statistics in m2 implementation

```

69 Generating CUDA Kernel Statistics...
70
71 Generating CUDA Memory Operation Statistics...
72 CUDA Kernel Statistics (nanoseconds)
73
74 Time(%)    Total Time    Instances    Average    Minimum    Maximum    Name
75 -----
76 100.0      72474164         2    36237082.0    14843675    57630489    conv_forward_kernel
77  0.0         2784         2      1392.0        1312         1472    do_not_remove_this_kernel
78  0.0         2624         2      1312.0        1280         1344    prefn_marker_kernel

```

Kernel statistics in optimization 1

Based on the result from nsys, the conv_forward_kernel total time of optimization 1 is 72.47 ms in total which is less than that of m2 implementation, 76.77 ms.

Kernel code

```
8  __global__ void conv_forward_kernel(float *output, const float *input, const float *mask, const int Batch,
9  const int Map_out, const int Channel, const int Height, const int Width, const int K)
10 {
11     /*
12     Modify this function to implement the forward pass described in Chapter 16.
13     We have added an additional dimension to the tensors to support an entire mini-batch
14     The goal here is to be correct AND fast.
15
16     Function paramter definitions:
17     output - output
18     input - input
19     mask - convolution kernel
20     Batch - batch_size (number of images in x)
21     Map_out - number of output feature maps
22     Channel - number of input feature maps
23     Height - input height dimension
24     Width - input width dimension
25     K - kernel height and width (K x K)
26     */
27     const int Height_out = Height - K + 1;
28     const int Width_out = Width - K + 1;
29     const int W_grid = ceil((Width_out*1.0)/TILE_WIDTHH);
30     // We have some nice #defs for you below to simplify indexing. Feel free to use them, or create your own.
31     // An example use of these macros:
32     // float a = in_4d(0,0,0,0)
33     // out_4d(0,0,0,0) = a
34
35     #define out_4d(i3, i2, i1, i0) output[(i3) * (Map_out * Height_out * Width_out) + (i2) * (Height_out * Width_out)
36     #define in_4d(i3, i2, i1, i0) input[(i3) * (Channel * Height * Width) + (i2) * (Height * Width) + (i1) * (Width)
37     #define mask_4d(i3, i2, i1, i0) kernel_mask[(i3) * (Channel * K * K) + (i2) * (K * K) + (i1) * (K) + i0]
38
39     // Insert your GPU convolution kernel code here
40
41     int n, m, h, w, c; // image_index, map_inedx, specific_height, specific_width, channel
42     n = blockIdx.x;
43     m = blockIdx.y;
44     h = (blockIdx.z / W_grid) * TILE_WIDTHH + threadIdx.y;
45     w = (blockIdx.z % W_grid) * TILE_WIDTHH + threadIdx.x;
46     float acc = 0;
47     if(h < Height_out && w < Width_out){
48         for(c = 0; c < Channel; ++c){
49             for(int p = 0; p < K; ++p){ // for loop, the mask K x K
50                 for(int q = 0; q < K; ++q){
51                     acc += in_4d(n, c, h+p, w+q) * mask_4d(m,c,p,q);
52                 }
53             }
54         }
55         out_4d(n,m,h,w) = acc;
56     }
57
58     #undef out_4d
59     #undef in_4d
60     #undef mask_4d
61 }
```

2. Optimization 2: tile shared memory convolution (2 points)

- a. Which optimization did you choose to implement? Chose from the optimization below by clicking on the check box and explain why did you choose that optimization technique.

- ☒ Tiled shared memory convolution (2 points)
- ☐ Shared memory matrix multiplication and input matrix unrolling (3 points)
- ☐ Kernel fusion for unrolling and matrix-multiplication (2 points)
- ☐ Weight matrix in constant memory (1 point)
- ☐ Tuning with restrict and loop unrolling (3 points)

- ☐ Sweeping various parameters to find best values (1 point)
- ☐ Multiple kernel implementations for different layer sizes (1 point)
- ☐ Input channel reduction: tree (3 point)
- ☐ Input channel reduction: atomics (2 point)
- ☐ Fixed point (FP16) arithmetic. (4 points)
- ☐ Using Streams to overlap computation with data transfer (4 points)
- ☐ An advanced matrix multiplication algorithm (5 points)
- ☐ Using Tensor Cores to speed up matrix multiplication (5 points)
- ☐ Overlap-Add method for FFT-based convolution (8 points)
- ☐ Other optimizations: please explain

I chose Tile shared memory convolution because it appeared on one of the previous MPs. In addition, previous lectures talked about this technique. I thought it was easy to implement.

- b. How does the optimization work? Did you think the optimization would increase performance of the forward convolution? Why? Does the optimization synergize with any of your previous optimizations?

Optimization 2 works by preloading data from global memory to shared memory. During the calculations of the convolution output, the threads can directly retrieve the data from shared memory rather than from the global memory. I thought this optimization would increase the performance of the forward convolution because shared memory is much faster than global memory. This can reduce global memory access and optimize memory utilization. This optimization synergizes with optimization 1: Weight Matrix in constant memory.

- c. List the Op Times, whole program execution time, and accuracy for batch size of 100, 1k, and 10k images using this optimization (including any previous optimizations also used).

Batch Size	Op Time 1	Op Time 2	Total Execution Time	Accuracy
100	0.164547 ms	0.657649 ms	0m1.244s	0.86
1000	1.55923 ms	6.53511 ms	0m9.735s	0.886
10000	15.415 ms	64.9177 ms	1m34.864s	0.8714

- d. Was implementing this optimization successful in improving performance? Why or why not? Include profiling results from *nsys* and *Nsight-Compute* to justify your answer, directly comparing to your baseline (or the previous optimization this one is built off of)

The optimization was not successful in improving performance.

```

54 Generating CUDA API Statistics...
55 CUDA API Statistics (nanoseconds)
56
57 Time(%)    Total Time    Calls    Average    Minimum    Maximum    Name
58 -----
59 78.8      1045156053      6      174192675.5      11977      559518420      cudaMemcpy
60 14.3      190061983       6      31676997.2       274319      186540130      cudaMalloc
61 5.5       72498446        6      12083074.3       2864       57633579      cudaDeviceSynchronize
62 1.2       16015083        6      2669180.5        22483      15879896      cudaLaunchKernel
63 0.2       2557808         6      426301.3         93130      797976       cudaFree
64 0.0       351282          2      175641.0         174328      176954       cudaMemcpyToSymbol
65
66
67
68
69 Generating CUDA Kernel Statistics...
70
71 Generating CUDA Memory Operation Statistics...
72 CUDA Kernel Statistics (nanoseconds)
73
74 Time(%)    Total Time    Instances    Average    Minimum    Maximum    Name
75 -----
76 100.0      72474164      2      36237082.0      14843675      57630489      conv_forward_kernel
77 0.0        2784         2      1392.0          1312         1472         do_not_remove_this_kernel
78 0.0        2624         2      1312.0          1280         1344         prefn_marker_kernel
79

```

Optimization 1 on baseline m2, nsys statistics

```

53 Generating CUDA API Statistics...
54 CUDA API Statistics (nanoseconds)
55
56 Time(%)    Total Time    Calls    Average    Minimum    Maximum    Name
57 -----
58 78.3      1031570868      6      171928478.0      17039      542470367      cudaMemcpy
59 14.1      186416526      6      31069421.0      264891      183215412      cudaMalloc
60 6.2       81202654       6      13533775.7       3107       65695059      cudaDeviceSynchronize
61 1.2       16101339       6      2683556.5       19550      15978940      cudaLaunchKernel
62 0.2       2439506        6      406584.3         87041      764002       cudaFree
63 0.0       351584         2      175792.0         174699      176885       cudaMemcpyToSymbol
64
65
66
67
68 Generating CUDA Kernel Statistics...
69
70 Generating CUDA Memory Operation Statistics...
71 CUDA Kernel Statistics (nanoseconds)
72
73 Time(%)    Total Time    Instances    Average    Minimum    Maximum    Name
74 -----
75 100.0      81181902      2      40590951.0      15487593      65694309      conv_forward_kernel
76 0.0        2912         2      1456.0          1408         1504         do_not_remove_this_kernel
77 0.0        2688         2      1344.0          1280         1408         prefn_marker_kernel
78

```

Optimization 1 & 2 on baseline m2, nsys statistics

Compared with optimization 1, optimization 1&2 spent less time on cudaMemcpy and cudaMalloc. However, the total time of conv_foward_kernel in optimization 1&2 is higher than that of conv_forward_kernel in optimization 1.

m3-const.ncu-rep X

Page: Details Launch: 1- 122 - conv_forward_kernel Add Baseline Apply Rules Copy as Image

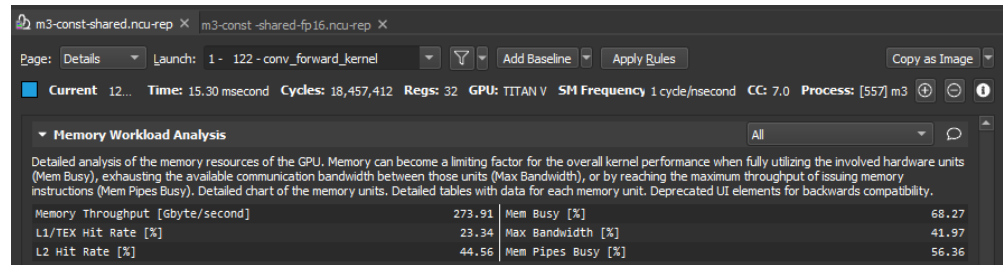
Current 12... Time: 14.62msecond Cycles: 17,540,190 Regs: 32 GPU: TITAN V SM Frequency 0 cycle/nsecond CC: 7.0 Process: [697] m3

Memory Workload Analysis

Detailed analysis of the memory resources of the GPU. Memory can become a limiting factor for the overall kernel performance when fully utilizing the involved hardware units (Mem Busy), exhausting the available communication bandwidth between those units (Max Bandwidth), or by reaching the maximum throughput of issuing memory instructions (Mem Pipes Busy). Detailed chart of the memory units. Detailed tables with data for each memory unit. Deprecated UI elements for backwards compatibility.

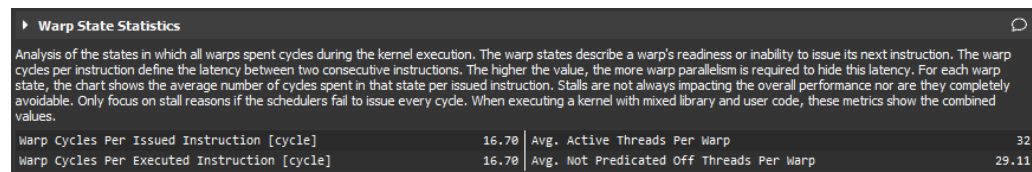
Memory Throughput [Gbyte/second]	286.69	Mem Busy [%]	71.37
L1/TEX Hit Rate [%]	95.19	Max Bandwidth [%]	47.36
L2 Hit Rate [%]	37.03	Mem Pipes Busy [%]	59.31

Optimization 1 on baseline m2, memory statistics

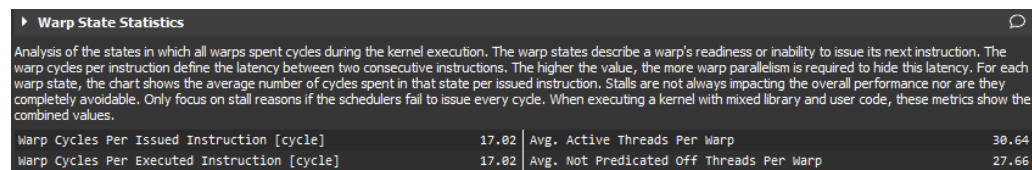


Optimization 1 & 2 on baseline: memory statistics

Compared with Optimization 1, the memory throughput in optimization 1&2 (273.91 GB/s) is lower than that of optimization 1 (286.69 GB/s). Even though optimization 2 has a higher L2 Hit Rate (44.56%), it has a lower L1 Hit Rate (23.34%). This is one factor that this optimization did not have a great performance.



Optimization 1 on baseline: warp statistics



Optimization 1&2 on baseline: warp statistics

Compared with optimization 1, optimization 1 & 2 has fewer average active threads per warp. Optimization 1 has an average of 32, and optimization 2 has an average of 30.64. The warp cycles per issued/executed instruction in optimization 2 are not as good as the ones in optimization 1. The poor performance can be caused by the kernel code that transfers data from global memory to shared memory. Not all threads in a warp would transfer the data to the shared memory, which introduced another set of control divergences to my kernel and lowered the number of active threads. This would also bring down the performance.

- e. What references did you use when implementing this technique?

MP4: 3D convolution

Lectures: lecture 7 and lecture 8

Using Shared Memory in CUDA C/C++: <https://developer.nvidia.com/blog/using-shared-memory-cuda-cc/>

- f. Please Paste your kernel code for this optimization. Your code should include the non-trivial code that you have changed for this optimization.

For example, it can be the complete kernel code for Tiled shared memory convolution several lines of code for Weight matrix in constant memory, or the “for” loop for loop unrolling

Code can be found under a folder called optimize in code submission.

Dynamically shared memory allocation, line 109.

```
93 __host__ void GPUInterface::conv_forward_gpu(float *device_output, const float *device_input, const float *device_mask, const int Batch, const int Map_out, const int Channel,
94 {
95     // Set the kernel dimensions and call the kernel
96     int H_out = Height - K + 1;
97     int W_out = Width - K + 1;
98
99     // defined the grids required for the output
100     int H_grid = ceil((H_out*1.0)/TILE_WIDTH); // height grids
101     int W_grid = ceil((W_out*1.0)/TILE_WIDTH); // width grids
102     int Z = H_grid * W_grid; // h x w requirements for the image
103     int blocksize = TILE_WIDTH + K - 1;
104
105     // dim3 DimBlocks(CACHE_SIZE,CACHE_SIZE, 1); // block dimension
106     dim3 DimBlocks(TILE_WIDTH,TILE_WIDTH, 1); // the un-optimized one
107     dim3 DimGrids(Batch, Map_out, Z); // grid dimension
108
109     conv_forward_kernel<<<DimGrids, DimBlocks, Channel * blocksize * blocksize * sizeof(float)>>>>(device_output, device_input, device_mask, Batch, Map_out, Channel, Height, Width, K);
110     // cudaDeviceSynchronize();
111 }
112 }
```

Kernel code:

```
#define TILE_WIDTH 16
__constant__ float kernel_mask[1 * 7 * 7 * 4 * 16];
__global__ void conv_forward_kernel(float *output, const float *input, const float *mask, const int
Batch, const int Map_out, const int Channel, const int Height, const int Width, const int K)
{
    const int Height_out = Height - K + 1;
    const int Width_out = Width - K + 1;
    const int W_grid = ceil((Width_out*1.0)/TILE_WIDTH);
    const int blocksize = TILE_WIDTH + K - 1;

    extern __shared__ float tileMem[]; // shared memory

    #define out_4d(i3, i2, i1, i0) output[(i3) * (Map_out * Height_out * Width_out) + (i2) *
(Height_out * Width_out) + (i1) * (Width_out) + i0]
    #define in_4d(i3, i2, i1, i0) input[(i3) * (Channel * Height * Width) + (i2) * (Height * Width) +
(i1) * (Width) + i0]
    #define mask_4d(i3, i2, i1, i0) kernel_mask[(i3) * (Channel * K * K) + (i2) * (K * K) + (i1) * (K)
+ i0]

    #define shareMem(i2, i1, i0) tileMem[(i2)*(blocksize*blocksize) + (i1) * blocksize + (i0)] // use
for 3d shared memory

    // Insert your GPU convolution kernel code here
    int tx = threadIdx.x;
    int ty = threadIdx.y;
```



```

int c; // image_index, map_inedx, specific_height, specific_width, channel
int h_topleft= (blockIdx.z / W_grid) * TILE_WIDTH;
int w_topleft = (blockIdx.z % W_grid) * TILE_WIDTH;
int n = blockIdx.x;
int m = blockIdx.y;
int h = h_topleft + ty; // the output height index
int w = w_topleft + tx; // the output width index
float acc = 0.0;

// with shared tiles, may be able to reduce sync_threads
for(c = 0; c < Channel; ++c){
    for(int i = ty; i < blocksize; i+= TILE_WIDTH){
        for(int j = tx; j < blocksize; j += TILE_WIDTH){
            if(h_topleft + i < Height && w_topleft + j < Width){
                shareMem(c,i,j) = in_4d(n, c, h_topleft + i, w_topleft + j);
            }
            else{
                shareMem(c,i,j) = 0.0f;
            }
        }
    }
}
__syncthreads();
if (h < Height_out && w < Width_out){
    for(c = 0; c < Channel; ++c){
        for(int p = 0; p < K; ++p){ // for loop, the mask K x K
            for(int q = 0; q < K; ++q){
                acc += shareMem(c ,ty+p , tx+q) * mask_4d(m,c,p,q);
                // acc += in_4d(n, c, h+p, w+q) * mask_4d(m,c,p,q);
            }
        }
    }
    out_4d(n,m,h,w) = acc;
}
#undef out_4d
#undef in_4d
#undef mask_4d
#undef shareMem
}

```

3. Optimization 3: Fixed point (FP16) arithmetic (4 points)

- a. Which optimization did you choose to implement? Chose from the optimization below by clicking on the check box and explain why did you choose that optimization technique.

- ☐ Tiled shared memory convolution (2 points)
- ☐ Shared memory matrix multiplication and input matrix unrolling (3 points)
- ☐ Kernel fusion for unrolling and matrix-multiplication (2 points)
- ☐ Weight matrix in constant memory (1 point)
- ☐ Tuning with restrict and loop unrolling (3 points)
- ☐ Sweeping various parameters to find best values (1 point)
- ☐ Multiple kernel implementations for different layer sizes (1 point)
- ☐ Input channel reduction: tree (3 point)
- ☐ Input channel reduction: atomics (2 point)
- ☒ Fixed point (FP16) arithmetic. (4 points)
- ☐ Using Streams to overlap computation with data transfer (4 points)
- ☐ An advanced matrix multiplication algorithm (5 points)
- ☐ Using Tensor Cores to speed up matrix multiplication (5 points)
- ☐ Overlap-Add method for FFT-based convolution (8 points)
- ☐ Other optimizations: please explain

I chose this optimization because I thought this optimization would be interesting since it would change the accuracy according to the final project GitHub. In addition, I found relevant documents on FP16 arithmetic. I thought that FP16 might have some advantage over float in calculation and size.

- b. How does the optimization work? Did you think the optimization would increase performance of the forward convolution? Why? Does the optimization synergize with any of your previous optimizations?

The optimization would convert the floating points number into a new type __half (FP16). Instead of using floating point arithmetic in calculating the convolution output, the kernel calculates the convolution output with FP16 arithmetic. In the end, the result will convert back to float. I think this will improve the performance of the forward convolution. FP16 is smaller than float, which means FP16 may have a faster computation speed due to the advantage of a smaller data size. Also, FP16 requires less memory than float. I thought this optimization synergizes with optimization 1.

- c. List the Op Times, whole program execution time, and accuracy for batch size of 100, 1k, and 10k images using this optimization (including any previous optimizations also used).

Batch Size	Op Time 1	Op Time 2	Total Execution Time	Accuracy
100	0.213373 ms	0.791051 ms	0m1.248s	0.86
1000	1.99508 ms	7.73802 ms	0m10.106s	0.887
10000	19.7247 ms	77.7143 ms	1m36.354s	0.8716

- d. Was implementing this optimization successful in improving performance? Why or why not? Include profiling results from *nsys* and *Nsight-Compute* to justify your answer, directly comparing to your baseline (or the previous optimization this one is built off)

This optimization was not successful in improving performance. This optimization was built on top of optimization 1, and the Op Time for optimization 1 is much better. The accuracy was slightly higher than the accuracy in the previous implementation. This was caused by the precision of calculations and rounding in FP16.

```

68 Generating CUDA Kernel Statistics...
69
70 Generating CUDA Memory Operation Statistics...
71 CUDA Kernel Statistics (nanoseconds)
72
73 Time(%)    Total Time    Instances    Average    Minimum    Maximum    Name
74 -----
75 100.0      98327679        2    49163839.5    19617856    78709823    conv_forward_kernel
76 0.0         2784          2     1392.0        1376        1408    do_not_remove_this_kernel
77 0.0         2688          2     1344.0        1312        1376    prefn_marker_kernel

```

Optimization 1 & 3 on baseline m2, took 98.33ms

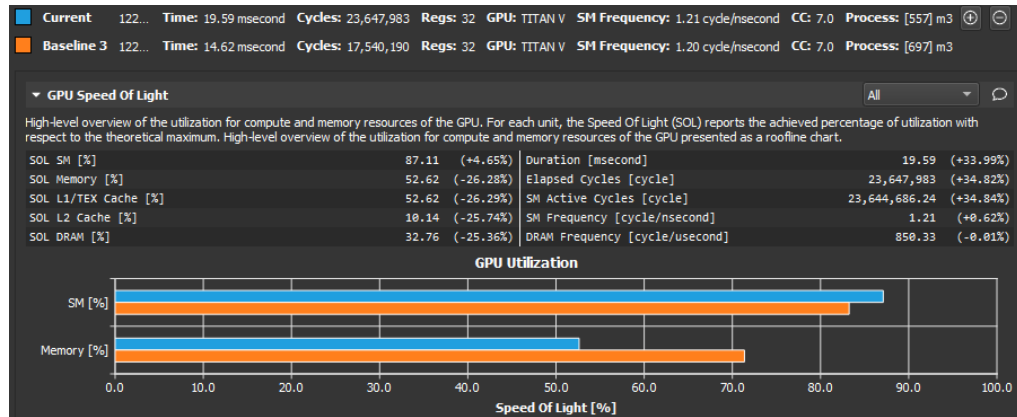
```

69 Generating CUDA Kernel Statistics...
70
71 Generating CUDA Memory Operation Statistics...
72 CUDA Kernel Statistics (nanoseconds)
73
74 Time(%)    Total Time    Instances    Average    Minimum    Maximum    Name
75 -----
76 100.0      72474164        2    36237082.0    14843675    57630489    conv_forward_kernel
77 0.0         2784          2     1392.0        1312        1472    do_not_remove_this_kernel
78 0.0         2624          2     1312.0        1280        1344    prefn_marker_kernel

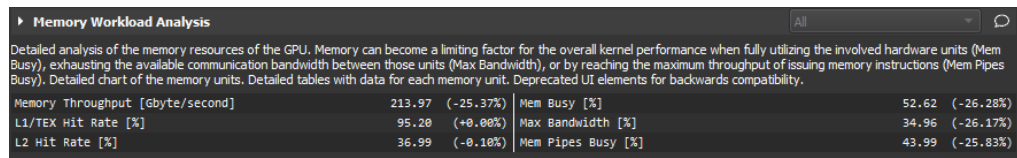
```

Optimization 1 on baseline m2, took 72.47ms

Based on the data from nsys, the time that conv_forward_kernel spent in optimization 1 & 3 is larger than that of optimization 1.

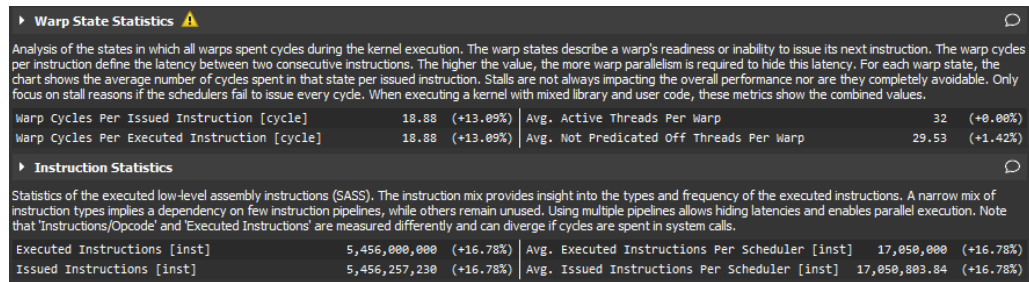


Blue: Optimization 1 & 3 on baseline m2, Orange: Optimization 1 on baseline m2



Memory statistics of Optimization 1&3 on baseline m2

Based on the result from the Nsight Compute, optimization 1&3 has a lower memory utilization and higher SM utilization. This indicates that optimization 1&3 does less memory in the kernel. However, the memory throughput of optimization 1&3 is 25.37% lower than the memory throughput of optimization 1.



Warp statistics and instruction statistics of Optimization 1&3 on baseline

The warp cycles per issue/execute instruction in optimization 1&3 are higher than that of optimization 1, so optimization 1&3 would have a longer instruction execution time. From the instruction statistics, the number of instructions in optimization 1&3 is 16.78% higher than the number of instructions in optimization 1. This is mainly caused by the conversion between float and FP16 (__half) data type. Although the FP16 data type may have some advantages over float type, the overheads of float to FP16 conversions and FP16 to float conversions cannot be avoided in optimization 1&3. Thus, the performance of optimization 1&3 is lower.

- e. What references did you use when implementing this technique?

CUDA Toolkit documentation:

<https://docs.nvidia.com/cuda/cuda-math-api/index.html>

- f. Please Paste your kernel code for this optimization. Your code should include the non-trivial code that you have changed for this optimization.
For example, it can be the complete kernel code for Tiled shared memory convolution several lines of code for Weight matrix in constant memory, or the “for” loop for loop unrolling

Code can be found under a folder called optimize in code submission.

Using a new library, cuda_fp16.h

```
1  #include <cmath>
2  #include <iostream>
3  #include <cuda_fp16.h>
4  #include "gpu-new-forward.h"
```

Kernel code

```
#define TILE_WIDTH 16
__constant__ float kernel_mask[1 * 7 * 7 * 4 * 16];

__global__ void conv_forward_kernel(float *output, const float *input, const
float *mask, const int Batch,
    const int Map_out, const int Channel, const int Height, const int Width,
const int K)
{
    const int Height_out = Height - K + 1;
    const int Width_out = Width - K + 1;
    const int W_grid = ceil((Width_out*1.0)/TILE_WIDTH);

    #define out_4d(i3, i2, i1, i0) output[(i3) * (Map_out * Height_out *
Width_out) + (i2) * (Height_out * Width_out) + (i1) * (Width_out) + i0]
    #define in_4d(i3, i2, i1, i0) input[(i3) * (Channel * Height * Width) +
(i2) * (Height * Width) + (i1) * (Width) + i0]
    #define mask_4d(i3, i2, i1, i0) kernel_mask[(i3) * (Channel * K * K) + (i2)
* (K * K) + (i1) * (K) + i0]

    // Insert your GPU convolution kernel code here

    int n, m, h, w, c; // image_index, map_inedx, specific_height,
specific_width, channel
    n = blockIdx.x;
    m = blockIdx.y;
    h = (blockIdx.z / W_grid) * TILE_WIDTH + threadIdx.y;
```

```

w = (blockIdx.z % W_grid) * TILE_WIDTH + threadIdx.x;
__half acc = __float2half_rn(0.0f);
if(h < Height_out && w < Width_out){
    for(c = 0; c < Channel; ++c){
        for(int p = 0; p < K; ++p){          // for loop, the mask K x K
            for(int q = 0; q < K; ++q){
                acc = __hadd(acc, __hmul(__float2half_rn(in_4d(n, c, h+p,
w+q)), __float2half_rn(mask_4d(m,c,p,q))));
            }
        }
    }
    out_4d(n,m,h,w) = __half2float(acc);
}

#undef out_4d
#undef in_4d
#undef mask_4d
}

```

4. Optimization 4: Using Streams to overlap computation with data transfer (4 point)

- a. Which optimization did you choose to implement? Chose from the optimization below by clicking on the check box and explain why did you choose that optimization technique.

- ☐ Tiled shared memory convolution (2 points)
- ☐ Shared memory matrix multiplication and input matrix unrolling (3 points)
- ☐ Kernel fusion for unrolling and matrix-multiplication (2 points)
- ☐ Weight matrix in constant memory (1 point)
- ☐ Tuning with restrict and loop unrolling (3 points)
- ☐ Sweeping various parameters to find best values (1 point)
- ☐ Multiple kernel implementations for different layer sizes (1 point)
- ☐ Input channel reduction: tree (3 point)
- ☐ Input channel reduction: atomics (2 point)
- ☐ Fixed point (FP16) arithmetic. (4 points)
- ☒ Using Streams to overlap computation with data transfer (4 points)
- ☐ An advanced matrix multiplication algorithm (5 points)
- ☐ Using Tensor Cores to speed up matrix multiplication (5 points)
- ☐ Overlap-Add method for FFT-based convolution (8 points)
- ☐ Other optimizations: please explain

I choose using streams to overlap computation with data transfer because there is a lecture talk about this. Implementing streams can increase the overall performance and data transfer rate.

- b. How does the optimization work? Did you think the optimization would increase performance of the forward convolution? Why? Does the optimization synergize with any of your previous optimizations?

This optimization works by creating multiple streams to overlap kernel execution and data transfer. A large chunk of data is broken into smaller chunks and distributed among the streams. I thought this would increase the overall performance. Some streams might execute the kernel whereas other streams transfer the data and prepare for kernel execution. This creates parallelism among the streams and reduces a certain amount of latency in data transfer. I think the optimization synergizes with optimization 1 and produces better performance.

- c. List the Op Times, whole program execution time, and accuracy for batch size of 100, 1k, and 10k images using this optimization (including any previous optimizations also used).

Batch Size	Op Time 1	Op Time 2	Total Execution Time	Accuracy
100	0.003278 ms	0.002823 ms	0m1.236s	0.86
1000	0.002481 ms	0.003544 ms	0m9.889s	0.886
10000	0.003199 ms	0.003446 ms	1m38.379s	0.8714

- d. Was implementing this optimization successful in improving performance? Why or why not? Include profiling results from *nsys* and *Nsight-Compute* to justify your answer, directly comparing to your baseline (or the previous optimization this one is built off of)

The optimization was not very successful in improving performance. Based on the nsys output data below, even though optimization 1&4 have better total time in memory operation (454.43 ms) than the total time in memory operation of optimization 1 (1040.49 ms), the total time in conv_forward_kernel of optimization 1&4 (86.61 ms) did not beat the total time of optimization 1 (72.47ms).

```

71 Generating CUDA Kernel Statistics...
72
73 Generating CUDA Memory Operation Statistics...
74 CUDA Kernel Statistics (nanoseconds)
75
76 Time(%)    Total Time    Instances    Average    Minimum    Maximum    Name
77 -----
78 100.0      86609572      800         108262.0   43007     190814    conv_forward_kernel
79 0.0        2816         2          1408.0     1376      1440     do_not_remove_this_kernel
80 0.0        2656         2          1328.0     1280      1376     prefn_marker_kernel
81
82
83 CUDA Memory Operation Statistics (nanoseconds)
84
85 Time(%)    Total Time    Operations    Average    Minimum    Maximum    Name
86 -----
87 73.5      334038815     800         417548.5   143007     764633    [CUDA memcpy DtoH]
88 26.5      120391699     806         149369.4   1504      38767478  [CUDA memcpy HtoD]

```

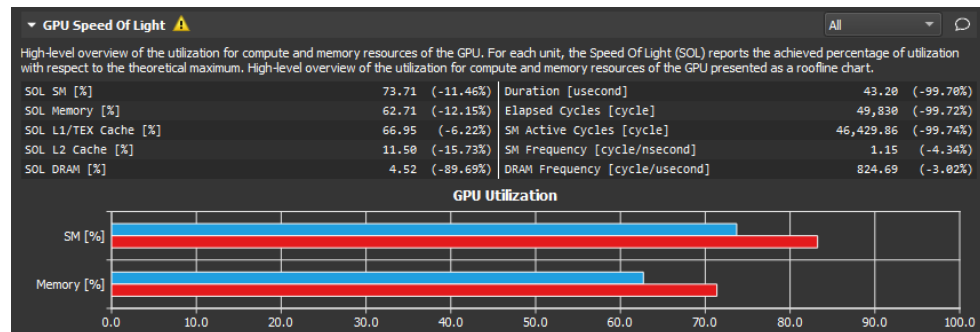
Optimization 1&4 on baseline m2

```

69 Generating CUDA Kernel Statistics...
70
71 Generating CUDA Memory Operation Statistics...
72 CUDA Kernel Statistics (nanoseconds)
73
74 Time(%)    Total Time    Instances    Average    Minimum    Maximum    Name
75 -----
76 100.0      72474164      2      36237082.0    14843675    57630489    conv_forward_kernel
77 0.0        2784        2      1392.0        1312        1472        do_not_remove_this_kernel
78 0.0        2624        2      1312.0        1280        1344        prefn_marker_kernel
79
80
81 CUDA Memory Operation Statistics (nanoseconds)
82
83 Time(%)    Total Time    Operations    Average    Minimum    Maximum    Name
84 -----
85 91.4      950719186      2      475359593.0    392019063    558700123    [CUDA memcpy DtoH]
86 8.6       89773144      6      14962190.7    1504        47988528    [CUDA memcpy HtoD]
87

```

Optimization 1 on baseline m2



Blue: optimization 1&4 on baseline m2, Red: optimization 1 on baseline m2

Memory Workload Analysis

Detailed analysis of the memory resources of the GPU. Memory can become a limiting factor for the overall kernel performance when fully utilizing the involved hardware units (Mem Busy), exhausting the available communication bandwidth between those units (Max Bandwidth), or by reaching the maximum throughput of issuing memory instructions (Mem Pipes Busy). Detailed chart of the memory units. Detailed tables with data for each memory unit. Depreciated UI elements for backwards compatibility.

Memory Unit	Value	Change	Memory Unit	Value	Change
Memory Throughput [Gbyte/second]	28.65	(-90.01%)	Mem Busy [%]	63.18	(-11.48%)
L1/TEX Hit Rate [%]	95.46	(+0.28%)	Max Bandwidth [%]	41.54	(-12.27%)
L2 Hit Rate [%]	94.17	(+154.29%)	Mem Pipes Busy [%]	53.06	(-10.53%)

Optimization 1&4 on baseline m2 compared with Optimization 1 on m2

Based on the information from Nsight Compute, both utilization on SM and memory in optimization decreases. The memory throughput of each kernel launch in optimization 1&4 is not as good as the one in optimization 1. The good things are that the L1 and L2 hit rate in optimization 1&4 are high. The decrease in the performance may come from the number of streams that I used in my code and data access patterns. The number of streams that I picked may not be optimized one for this convolution kernel. The amount of overlapping between kernel execution and data transfer is not sufficient.

- e. What references did you use when implementing this technique?

Lecture 22

- f. Please Paste your kernel code for this optimization. Your code should include the non-trivial code that you have changed for this optimization.

For example, it can be the complete kernel code for Tiled shared memory convolution several lines of code for Weight matrix in constant memory, or the “for” loop for loop unrolling

Code can be found under a folder called optimize in code submission.

Kernel code

```
#define TILE_WIDTH 16
#define STREAM_NUM 4
__constant__ float kernel_mask[1 * 7 * 7 * 4 * 16];

__global__ void conv_forward_kernel(float *output, const float *input, const
float *mask, const int Batch,
    const int Map_out, const int Channel, const int Height, const int Width,
const int K)
{
    const int Height_out = Height - K + 1;
    const int Width_out = Width - K + 1;
    const int W_grid = ceil((Width_out*1.0)/TILE_WIDTH);

    #define out_4d(i3, i2, i1, i0) output[(i3) * (Map_out * Height_out *
Width_out) + (i2) * (Height_out * Width_out) + (i1) * (Width_out) + i0]
    #define in_4d(i3, i2, i1, i0) input[(i3) * (Channel * Height * Width) +
(i2) * (Height * Width) + (i1) * (Width) + i0]
    #define mask_4d(i3, i2, i1, i0) kernel_mask[(i3) * (Channel * K * K) + (i2)
* (K * K) + (i1) * (K) + i0]

    // Insert your GPU convolution kernel code here

    int n, m, h, w, c; // image_index, map_inedx, specific_height,
specific_width, channel
    n = blockIdx.x;
    m = blockIdx.y;
    h = (blockIdx.z / W_grid) * TILE_WIDTH + threadIdx.y;
    w = (blockIdx.z % W_grid) * TILE_WIDTH + threadIdx.x;
    float acc = 0;
```

```

        if(h < Height_out && w < Width_out){
            for(c = 0; c < Channel; ++c){
                for(int p = 0; p < K; ++p){          // for loop, the mask K x K
                    for(int q = 0; q < K; ++q){
                        acc += in_4d(n, c, h+p, w+q) * mask_4d(m,c,p,q);
                    }
                }
            }
            out_4d(n,m,h,w) = acc;
        }

#undef out_4d
#undef in_4d
#undef mask_4d
}

```

create streams and declare some useful variables

```

65 __host__ void GPUInterface::conv_forward_gpu_prolog(const float *host_output, const float *host_input, const float *host_mask, float **device_output_ptr, float
66 {
67     // Allocate memory and copy over the relevant data structures to the GPU
68
69     // We pass double pointers for you to initialize the relevant device pointers,
70     // which are passed to the other two functions.
71
72     // Useful snippet for error checking
73     // cudaError_t error = cudaGetLastError();
74     // if(error != cudaSuccess)
75     // {
76     //     std::cout<<"CUDA error: "<<cudaGetErrorString(error)<<std::endl;
77     //     exit(-1);
78     // }
79     int H_out = Height - K + 1; // compute the output height
80     int W_out = Width - K + 1; // compute the input width
81     cudaMalloc((void**)device_output_ptr, Batch * Map_out * H_out * W_out * sizeof(float)); // using the output height and width
82     cudaMalloc((void**)device_input_ptr, Batch * Channel * Height * Width * sizeof(float)); // using the input height and width
83     cudaMemcpy(device_input_ptr, host_input, Batch * Channel * Height * Width * sizeof(float), cudaMemcpyHostToDevice); // copy host input to device input
84     cudaMemcpyToSymbol(kernel_mask, host_mask, 1 * 7 * 7 * 4 * 16*sizeof(float)); // copy to the constant memory
85
86
87     // defined the grids required for the output
88     int H_grid = ceil((H_out*1.0)/TILE_WIDTH); // height grids
89     int W_grid = ceil((W_out*1.0)/TILE_WIDTH); // width grids
90     int Z = H_grid * W_grid; // h x w requirements for the image
91     const int segSize = 25; // define seg size
92
93     cudaStream_t streams[STREAM_NUM];
94     for(int i = 0; i < STREAM_NUM; ++i){ // create streams
95         cudaStreamCreate(&streams[i]);
96     }
97
98     int input_CHW = Channel * Height * Width;
99     int output_MHW = Map_out * H_out * W_out;
100     int in_copy_size = segSize * input_CHW;
101     int out_copy_size = segSize * output_MHW;
102
103     dim3 DimBlocks(TILE_WIDTH, TILE_WIDTH, 1); // the un-optimized one
104     dim3 DimGrids(segSize, Map_out, Z); // grid dimension

```

The for loop for data transfer and kernel launch with streams

```
105
106 for (int i = 0; i < Batch; i += (STREAM_NUM * segSize)){
107
108     // input offsets
109     int offset0 = (i + 0 * segSize) * input_CHW;
110     int offset1 = (i + 1 * segSize) * input_CHW;
111     int offset2 = (i + 2 * segSize) * input_CHW;
112     int offset3 = (i + 3 * segSize) * input_CHW;
113
114     // output offset
115     int out_offset0 = (i + 0 * segSize) * output_MHW;
116     int out_offset1 = (i + 1 * segSize) * output_MHW;
117     int out_offset2 = (i + 2 * segSize) * output_MHW;
118     int out_offset3 = (i + 3 * segSize) * output_MHW;
119
120     // async cpy input, host to device
121     cudaMemcpyAsync(*device_input_ptr + offset0, host_input + offset0, in_copy_size * sizeof(float), cudaMemcpyHostToDevice, streams[0]);
122     cudaMemcpyAsync(*device_input_ptr + offset1, host_input + offset1, in_copy_size * sizeof(float), cudaMemcpyHostToDevice, streams[1]);
123     cudaMemcpyAsync(*device_input_ptr + offset2, host_input + offset2, in_copy_size * sizeof(float), cudaMemcpyHostToDevice, streams[2]);
124     cudaMemcpyAsync(*device_input_ptr + offset3, host_input + offset3, in_copy_size * sizeof(float), cudaMemcpyHostToDevice, streams[3]);
125
126     // stream kernel calls
127     conv_forward_kernel<<<DimGrids, DimBlocks, 0, streams[0]>>>(*device_output_ptr + out_offset0, *device_input_ptr + offset0, *device_mask_ptr, Batch, Map_out, Channel, Height, Width, K);
128     conv_forward_kernel<<<DimGrids, DimBlocks, 0, streams[1]>>>(*device_output_ptr + out_offset1, *device_input_ptr + offset1, *device_mask_ptr, Batch, Map_out, Channel, Height, Width, K);
129     conv_forward_kernel<<<DimGrids, DimBlocks, 0, streams[2]>>>(*device_output_ptr + out_offset2, *device_input_ptr + offset2, *device_mask_ptr, Batch, Map_out, Channel, Height, Width, K);
130     conv_forward_kernel<<<DimGrids, DimBlocks, 0, streams[3]>>>(*device_output_ptr + out_offset3, *device_input_ptr + offset3, *device_mask_ptr, Batch, Map_out, Channel, Height, Width, K);
131
132     // async cpy output, device to host
133     cudaMemcpyAsync((void*)(host_output + out_offset0), *device_output_ptr + out_offset0, out_copy_size * sizeof(float), cudaMemcpyDeviceToHost, streams[0]);
134     cudaMemcpyAsync((void*)(host_output + out_offset1), *device_output_ptr + out_offset1, out_copy_size * sizeof(float), cudaMemcpyDeviceToHost, streams[1]);
135     cudaMemcpyAsync((void*)(host_output + out_offset2), *device_output_ptr + out_offset2, out_copy_size * sizeof(float), cudaMemcpyDeviceToHost, streams[2]);
136     cudaMemcpyAsync((void*)(host_output + out_offset3), *device_output_ptr + out_offset3, out_copy_size * sizeof(float), cudaMemcpyDeviceToHost, streams[3]);
137 }
138
139 cudaFree(*device_input_ptr);
140 cudaFree(*device_output_ptr);
141 }
```

Other changed part of the code

```
144 __host__ void GPUInterface::conv_forward_gpu(float *device_output, const float *device_input, con
145 {
146     // Set the kernel dimensions and call the kernel
147     // int H_out = Height - K + 1;
148     // int W_out = Width - K + 1;
149
150     // // defined the grids required for the output
151     // int H_grid = ceil((H_out*1.0)/TILE_WIDTH); // height grids
152     // int W_grid = ceil((W_out*1.0)/TILE_WIDTH); // width grids
153     // int Z = H_grid * W_grid; // h x w requirements for the image
154
155
156     return;
157     // cudaDeviceSynchronize();
158 }
159
160
161
162 __host__ void GPUInterface::conv_forward_gpu_epilog(float *host_output, float *device_output, flo
163 {
164     // Copy the output back to host
165     // int H_out = Height - K + 1;
166     // int W_out = Width - K + 1;
167     // Free device memory
168     return;
169 }
```